

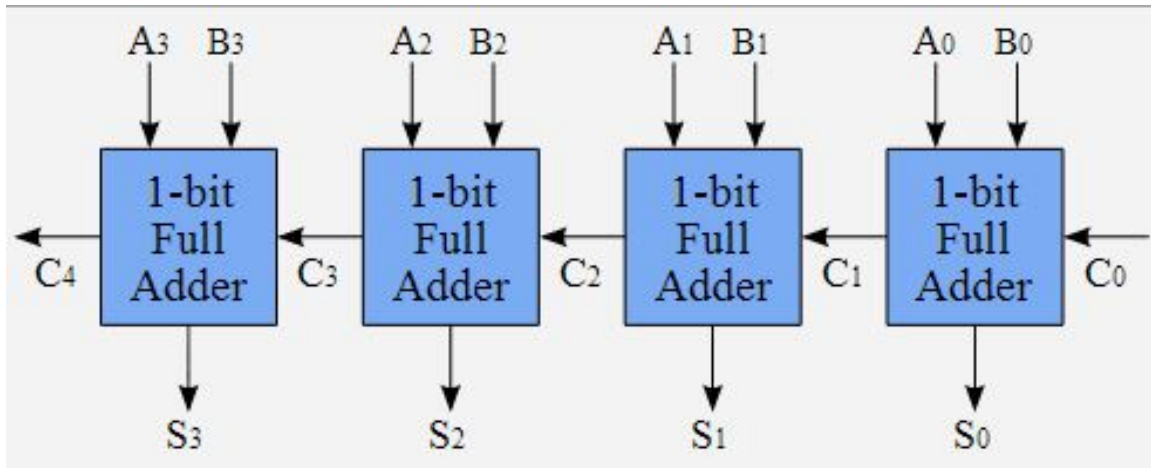
## Lecture 2

- CHDL or HDL is a textual language for designing hardware circuits.
- Circuits can be specified via a diagram schematic or a textual CHDL.
- Diagrams are easier for the beginner but not very scalable or useful.
- CDHLs offer better scalability and support formal methods.
- Circuits are very similar to functions. Black box, internal computation, produce outputs.
- Give circuits a name in Hydra. This type of 'black box' definition is similar to abstraction in imperative programming languages.
- Multiplexor in primitives:  $\text{mux1 } c \ a \ b = \text{or2 } (\text{and2 } (\text{inv } c) \ a) \ (\text{and2 } c \ b)$ .
- Think of the multiplexor order as counting from 0 to highest ed  $\text{mux2 } (0,1) \ p \ q \ r \ s = q$  (second value).
- Demultiplexors take an control input and a data input and produce 2 outputs. The control selects which output to put the data on.
- Demux2 and Mux2 can be built by combining demux1s or mux1s
- Tuples used where there is no logical grouping and indexing doesn't make sense.
- Words used when there is a logical sense of numbering and ordering.
- Words of Tuples and Tuples or Words are allowed.

## Lecture 3

- No feedback loops in combinational logic.
- Make sure that inputs and outputs are known BEFORE implementation.
- Efficiency is tied in with a circuits cost model. Not necessarily the number of components.
- It is generally better to come up with a better algorithm.
- In VLSI regularity is better than random optimization.
  1. Synthesis from truth table.
  2. Synthesis from algebraic logic
  3. "Sum of cases" technique.
  4. Conditionals and cases.
  5. Using encodings.
  6. Defining building blocks
  7. Standard patterns
- "Sum of cases" technique is a way of truth table synthesis. (mapping inputs to outputs)
- Nested multiplexors are used for conditionals and cases.
- Cases correspond to  $2^k$  cases given a k-bit opcode. Use demultiplexors in these circumstances.

- Half adder - A sum and a carry. Full adder a 2 bit sum and a carry.
- Ripple adders are combinations of the 'full adder' building block




- Efficiency should always be done in reference to cost model.
- Optimization only plays a part later, it should not interfere with synthesis or it could become counterproductive.
- Cost models place a cost on a circuit. The best circuit is the one with the lowest cost.
- Examples of cost models include number of components (not useful nowadays as we have integrated circuits) and path depth (for combinational circuits).
- For k inputs. There are  $2^k$  lines in its truth table and  $2^{(2^k)}$  boolean functions of k inputs.
- General logic gates (LUT). Important in reconfigurable circuits.  $p = \text{lut2}(a,b,c,d) \times y$
- Gate delay is the time it takes for a signal to become valid once it has reached the gate.
- Difference between invalidity and incorrectness.
- For each signal x in the circuit, the time at which x becomes stable is called  $t(x)$ . This is the path depth of x
- Inputs to x are denoted  $I(x)$ .
- Gate delay of a gate = gate delay of previous gates and the gates own delay.

$$t(x) = d + \max_{y \in I(x)} t(y)$$

- Critical path depth is the maximum path depth of all of its outputs. This can be used to determine the 'speed' of the circuit.

## Lecture 4

- Hazards, incorrect output based on varying gate delays (ie  $\text{and2 } x (\text{inv } x)$  results in a 1 value while  $(\text{inv } x)$  suffers gate delay).
- A sequential circuit may have flip flops and feedback loops.
- Sequential circuits perform computations through time.
- Flip-flop is basic memory (hydra `def - dff :: Clocked a → a → a`)
- Clocked used a Bool is simply not rich enough.
- Real systems have a reset mechanism to put flip-flops into a valid state.
- System must ensure a valid flip-flop value (no mid-voltages).
- Flip-flops update on clock tick. Output continuous but only updates when clock pulses 1.
- Abstract clock and Physical clock differences. Abstract cannot exist in real life. Components treat a 0 to 1 transition as a tick.
- Synchronous model is used to prevent hazards. It boasts simple behavior and satisfies some constraints.
- Synchronous if:
  1. Contains logic gates and flip flops.
  2. Every flip flop is connected to a global clock.
  3. Every clock tick reaches each flip flop at the same time.
  4. Each feedback loop goes via a flip flop.
  5. Inputs to a circuit remain stable during a clock cycle.
- The restrictions on this mean that you cannot do combinational logic on clock input, you cannot introduce state via combinational feedback loops, and the circuit may malfunction if inputs are not stable for long enough.
- State can be created by using flip flops or feedback loops in combinational logic (difficult to reason).
- A synchronous circuit is a sequential circuit that behaves like a finite state automation.
- Synchronous trade off. Restricted design choice but much easier to reason the circuit.
- Clock is an implicit input (as all dffs have the same clock ALWAYS).
- Clock tick is a point in time. Clock cycle is an interval.
- Clock skew is when a clock tick does not reach flip flops at the same time (such as in a fanout or if there are longer wires etc).
- The floor plan of the clock is normally laid out well in advance. See example H-tree. Known as the clock distribution.



- Circuits can have 'global' state in the synchronous model (values of all flip flops).
- Circuit behaves as a single finite state automation.
- Clock removes concern for intermediate values, hazards and gate delays (as they are all dealt with in the synchronous model).
- reg1 component like a dff but with option ld signal that loads/ignores input value after clock tick accordingly.
- reg1 works by multiplexing the input to a dff depending on the value of ld (obviously involves a feedback loop).

## Lecture 5

- list structure for words similar to haskell notation ie.  $(x:xs) = [a,b,c,d,e]$ ,  $x = a$  and  $b = [b,c,d,e]$
- equally the ++, !! and length operators can be used like in haskell.
- field w i k takes the value of w starting at position i for k elements.
- The type of a circuit describes its input. Such as how it's represented, how many inputs/outputs it takes/produces, parameter size and building blocks (if any).
- You may use circuit generators to save making every component by hand. Eg. rippleAdd, a generic n-bit ripple adder that works for any word size.
- This is not software to do the operation, it is still a real circuit.
- Hydra has 4 'times'

### 1. Model Transformation Time

2. Circuit Generation Time
3. Simulator Compilation Time
4. Runtime

- Always use type a rather than Bool, its a hydra thing...
- bitslice2 can be used so that  $[a] \rightarrow [a] \rightarrow [(a,a)]$
- mux1w is like a multiplexor but outputs words rather than single bits.
- general n-bit reg an wlatch (wlatch doesn't take a load signal).
- Modern RISC and CISC architectures use an indexed set of registers R0, R1, R2 etc which are implemented in a circuit called the register file.
- Memory locations are usually bytes while registers contain words. Registers are few in comparison to memory (but much faster).
- A dual-port register file provides separate input from output. Some register files have three ports (one in and 2 out).
- How to handle address inputs. A tree of multiplexors and demultiplexors is needed for the addresses.
- Look over register file recursive definitions.

## Lecture 6

- A test bench takes readable input and produces readable output. It converts to and from the actual signals and formats the output.
- Lots of hydra test bench definitions. 'getbit', 'getbin' and 'gettc'
- Additional hydra functions 'fanout', 'shl', 'bitslice2', 'orw'
- An ALU is a combinational circuit. It is typically an adder with augmented functions to support other functions (eg bit shifting, subtraction, incrementing).
- The critical path of the ALU should be very small.
- Multiplication, division etc are performed over a series of clock ticks (by using sequential algorithms).
- Sequential functions may operate as a machine language program, using normal processing registers or using a dedicated sequential circuit with its own registers (functional unit).
- A functional unit is a sequential black box circuit with its own specific registers to perform a task like multiplication. It receives an input start and outputs an output signal 'ready'.
- By using a functional unit (which is sequential), we can accept it may take a few clock cycles which means the global clock doesn't have to be slowed down.
- If it were done by combinational logic the path depth would be so great the clock speed would have to be slowed down.
- Use the shift and add algorithm to do multiplication. The inputs to  $x*y$  would be x y and start, the outputs would be result and ready.

## Lecture 7

- Number of patterns in combinational logic; sum of cases, truth tables, fold, scan, map.
- Design patterns are higher order functions. They take one or more circuit specifications as parameters.
- The pattern defines how to connect up these given circuits in a regular pattern.
- Pattern definition similar to ordinary circuit specification except that: It uses recursion to decompose groups of signals, it uses abstract circuits, rather than specific circuits, its type may include building block circuits.
- map, map2, mapn used as in haskell.
- fold as well. and2/or2 operate by using left folds.
- ascanr is like fold but returns a list (including all intermediate values)

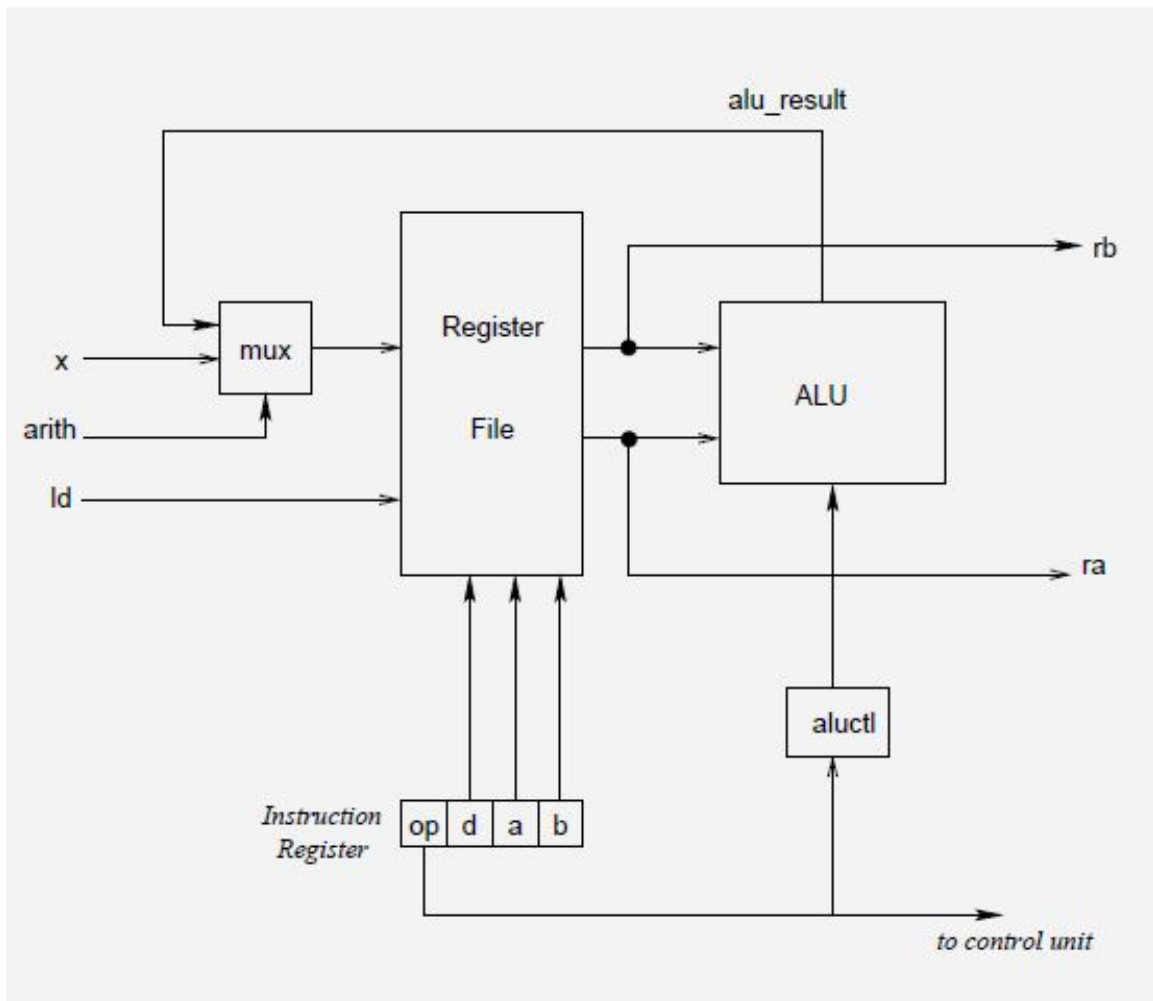
## Lecture 8

- FSA (finite state automation) has a set of states, a stream of inputs (and state conversion function). It is used to recognize if the grammar of an input is correct.
- FSAs need to be more general for practical applications. For example the new state transition function could produce a new output as well as defining a new state to enter.
- In an unencoded representation you could use a dff for each state. If exactly one flip flop contains a 1 then it is in that state.
- In an unencoded representation the no. of states = no. state flip flops.
- In an encoded representation, each possible state is represented by a unique binary number held in a register.
- In an encoded representation, no. of flip-flips =  $\log_2(\text{no. of states})$
- For the shift and add algorithm we must distinguish between what to put in the registers when start = 1 AND what to put in the registers when start  $\neq$  1 but a multiplication is still taking place.
- When start = 1 x & rx are both k bits, y i k bits, ry is 2k bits (it needs padding) and prod is 2k bits. rx and ry are set to the inputs on signals x & y and prod is set to 0.
- rx and ry are shifted right and left accordingly so that the algorithm is always examining the same bits.

## Lecture 9

- Sigma 16 is a 16-bit RISC style architecture. No bytes/words. Everything is 16-bit. 16 registers. Based on a simplified version of MIPS.
- RRR format (two operands in the register file. The result goes into a register)
- RX format (one operand is a register, the other is a memory address).

- Memory addresses are constants, they are the combination of the displacement and the index/base.
- RRR takes 4 inputs. The opcode op, destination d, and source addresses sa & sb.
- RX takes 4 inputs. The RX opcode (MUST be f), a destination register d, an index register sa and the secondary opcode sb.
- RRR (one word) RX (two words, 4 4 bit fields and a 16-bit displacement).
- Instructions can often be implemented without a feedback loop.



- The datapath contains the combinational logic for calculations.
- The control generates all of the control signals needed for the datapath
- Similar design complexities, the datapath focuses more on structure whereas the control focuses on behavior.
- Different design techniques (datapath designed directly as a circuit, control as an algorithm). The separation helps clarify timing.

- Design of the datapath (all the logic and registers needed for calculations) will logically create a set of control signals that operate it.
- There is a naming convention for control signals. ctl, subsystem, specific. Eg `ctl_rf_ld`.
- The datapath contains registers, computational systems and interconnections.
- It has two inputs, the control signals and a data word (from either memory system or the DMA input controller).
- It outputs the bus to memory an I/O (it also outputs the states of the registers).
- The datapath can be broken down into further subsystems (registers, ALU, busses etc).
- Introduce an instruction register and program counter. For RX instructions an `adr` register is needed for the second word

## Lecture 10

- `adr`, `pc` and `ir` needed in the datapath to keep track of program execution. ALU needed to perform arithmetic.
- The datapath contains the processor's registers.
- Aim to implement instruction set by an algorithm that uses the registers as temporary variables.
- Computations on registers are performed via combination logic in the datapath.
- registers and multiplexors have load control/control signals. The overall behavior is determined by these control signals.
- many datapath subsystems take control input. The right values must be place on all control signals to make the datapath do something useful.
- To perform an operation control signals and data inputs must be set (`memdat` coming from the system memory bus). Outputs will appear at the next clock cycle.
- almost always will the `ir` need to be set and the `pc` incremented. `ir := mem[pc]`, `pc := pc + 1` (the comma denotes that they occur in parallel).
- **\*\*Perhaps look at example operations\*\*** Lecture very practical based.

## Lecture 11

- General behavior of the datapath is defined by the control algorithm. Which is described with high level statements such as `ir := mem[pc]`.
- A circuit must be designed that will generate the correct control signals to be useful.
- High level control algorithm  $\rightarrow$  determine which signals must be asserted to achieve each statement  $\rightarrow$  synthesized control circuit to set the control signals.
- **\*\*Learn some examples of high level control RRR and RX functions\*\***
- Imperative languages have:




- Computational Statements - assignments etc.
  - Control Structures - concerned with what and when (eg if-then-else, while loops etc).
- A control language needs:
  - Computational Statement - which control signals should be set. (eg. Assert - which control signals are on/off).
  - Control Structures (similar to those in imperative languages).
- Straight-line control mode. A sequence of assert statements dictate what the control signals will be for a number of clock cycles.
- Statements can also be given names such as state\_1 : Assert [ctl\_a, ctl\_b, ctl\_c]. Meaning state state\_1 is synonymous with ctl\_a,b,c all being set to 1 and all others being set to 0.
- Control structures must also be considered, the most important are:
  - Repeat Forever
  - It-Then-Else
  - Case (multi-way conditional)
  - While
- Sequences of statements can be embedded in a 'block' (between begin and end). While they still require several clock cycles to execute, it can be convenient for grouping.
- Almost every control structure **REQUIRES** an infinite loop, otherwise the datapath would stop doing things as the control would terminate giving it control signals.
- fairly obvious case and if-then-else notation.
- From high level to low level. high level control algorithm can be used as an emulator (executes machine language at register transfer level).
- Aim is to translate this high level control algorithm into a circuit
  1. For each step, determine the control signals that must be set during a clock cycle.
  2. Develop a circuit that generates those signals.
- **\*\*More examples to look at\*\***
- Specify control algorithm, synthesize a control circuit from the algorithm, when changes are needed, modify the algorithm and re-synthesize. (part of synthesis of control).
- Synthesis techniques. The delay element method. Simple, flexible and efficient but complex and irregular wiring.
- Done by having 1 flip-flop per state. At any time only 1 flip-flop contains 1. If a flip-flop is in state 1 it means that corresponding statement is being executed.

```

state_1 : assert [ctl_b]
state_2 : assert [ctl_a, ctl_c]
state_3 : assert [ctl_a, ctl_b]
state_4 : assert [ctl_c]

```




[h]

## Generating the control signals

```

state_1 : assert [ctl_b]
state_2 : assert [ctl_a, ctl_c]
state_3 : assert [ctl_a, ctl_b]
state_4 : assert [ctl_c]

```



[h]

- \*\*Again more examples\*\*

## Lecture 12

- M1 is the first version of a circuit that implements the Sigma16 instruction set architecture. 2 ways to run Sigma16 programs, via an emulator or a circuit simulator on the M1 circuit.
- Several approaches to implement I/O in real systems.
  - I/O instructions, read and write. Only in early computers. V. Slow.
  - Start I/O instruction. Device generates interrupt when ready.
  - Dedicated I/O processor, running concurrently with CPU. Used in many supercomputers.
  - Direct Memory Access. I/O implemented by stores and fetches at 'magic' memory addresses.

- M1 doesn't support general I/O. It has a special facility for loading a machine language program into memory (DMA).
- DMA = 1, one word is place in memory each clock cycle, simulation driver is given a list of words and puts it in memory (one per cycle).
- When finished the simulation driver puts 1 on reset of one cycle which starts the control algorithm.
- **\*\*DMA examples\*\***
- Software tools include: Assembler, Emulator, Circuit, Circuit Simulator, Simulation Driver.
- Two main environments higher level (emulation) and lower level (circuit simulation).
- Higher level: a suite of tools (with GUI) including assembly and emulator. Done in Hydra. Good for assembling programs, referencing behavior and long executions.
- Lower level: Compile a Sigma16 program with M1 circuit, M1 simulation driver and Hydra library. Produces a standalone application that executes on the Sigma16 program by running it on the circuit.
- **\*\*Assessed Exercise Details and ArrayMax.hs M1 program\*\***
- M1 and subsystems output internal registers and signals so they can be observed during execution.
- Simulation driver is different from the ones seen so far as it supports DMA, input is a program (not raw signals) and detects significant output events.
- **\*\*More Examples\*\***
- IMPORTANT NOTE. The last 2/3 lectures have been rife with examples. These notes may not be sufficient to understand (and more importantly apply) the content described.

## Lecture 13

- Latency is the time required for an operation. There is memory, instruction and communication latency.
- Throughput is the amount of work performed per unit of time (eg. memory accesses per second).
- For overall performance throughput is important.
- For (re)solving bottlenecks, latency is often important.
- Amdahl's law observes that the total speed up of the system may be less than the speed up of a part of the system depending on how much time is spend in said part of the system.
- Example of Amdahl's law. If part A of a system can be sped up from 30s to 15s and is 30% of the system process (the other part takes 70s) then the increase is 85s/100s which is 1.17 overall speed up factor.
- Moore's law observes manufacturing trends. It states that the number of transistors per chip will double every two years (or some other number of years).

- Moore's law is based on the fact that integrated circuit fabrication technology makes gradual improvements which allows more components to be placed on chips.
- Moore's law is not about speed. It is about chip density. The only way to increase speed now is to add more processing cores to a chip.
- Performance should only really be measured in terms of its 'wall clock' time to run a piece of software.

$$T = ni \cdot cpi \cdot spc$$

*or*

$$T = \frac{ni \cdot cpi}{cps}$$

- Where  $T$  = execution time in seconds,  $ni$  = no. instructions in order to execute the program,  $cpi$  = average no. clock cycles per instruction,  $spc$  = seconds per clock cycle,  $cps$  = cycles per second.
- Reducing one of these factors does not necessarily reduce  $T$ .
- In practice, the critical path will be the path from a register in a register file, through the ALU and back to the register file.
- The M1 control algorithm will place whatever value is on the memory bus into the destination register. Even if the value is incorrect (for instance if it takes several clock cycles to compute).
- More sophisticated solutions include using pipelining stages, interleaved memory and super-scalar memory.
- The M1 circuit has no interrupts, memory protection or virtual memory. No concurrency. Long-running instructions must execute until completion while the rest of the processor waits.
- Pipelining is the application of parallelism to the steps in executing an instruction.
- General approaches to improving speed include: improving clock speed (reduce critical path), reduce no. clock cycles (pipelining and superscalar parallelism), improve speed of memory (parallelism by interleaved memory, exploit locality).
- The basic application of pipelining is to execute streams of instructions.
- 'overlapping the steps in consecutive multiplications'.
- In multiplication, it is the idea of unrolling iterations.
- The functional unit has one register for each of the variables, each clock cycle the new variable overwrites the old variable. This is modified to make a sequence of register triples, each cycle the variables move from one set of registers to the next.
- 2 related techniques for improving speed. RETIMING - introduce extra registers to break up long paths through combinational logic. Allows for faster clock.
- PIPELINING - a stream of inputs goes through a sequence of stages. Introduce parallel execution among the stages. Same kind of speed up as an assembly line in manufacturing.
- The idea of retiming is to find the critical path and break it into smaller paths with registers in between. It will not take several clock cycles to perform the critical path operation but the clock can now be made faster.

- The critical operation becomes slower due to the safety factor, whereas non-critical operations will benefit speed up.

## Lecture 14 (Incomplete)

- More stages in pipelining for RX instructions. (fetch second word of instruction, calculate ea, perform memory access).
- All instructions have to go through the same sequence of stages, so pipeline must have enough stages for the worst case.
- Conflicts caused by sequencing in pipelining are called hazards. Some can be circumvented by introducing more circuits, some require stalling the pipeline, so we get sequential execution for a brief time.
- Important to maintain correct behavior.
- To avoid structural hazards, we separate memories for instructions and data.
- Data hazards occur when the 'wrong' value is accessed. Similar to critical sections in imperative programming.
- There are 2 main classifications:
  1. Read After Write
  2. Write After Read
  3. Write After Write
  4. Intuitively, there is no hazard invoked from a Read After Read
- To detect hazards, we must make a complete list of hazards and introduce combinational logic to check them all. It is essential to find ALL the potential hazards.
- To detect a read after write hazard:

$$raw = (ir1\_sa = ir3\_d) \vee (ir\_sb = ir3\_d)$$

- To handle data hazards we can:
  1. STALLING - temporarily stop an instruction in an earlier stage from proceeding, giving time for a later stage to complete its work. This introduces a 'bubble' into the system.
  2. BYPASSING (also called forwarding) - introduce additional circuitry that allows the pipeline to continue at full speed without error.
- BYPASSING in summary, is when an early stage requires a value that hasn't yet been loaded into the register (but the value exists elsewhere in the machine). We can introduce a new signal that 'bypasses' the register and transmits the signal to where it is needed directly. More comparators and multiplexors.
- A functional unit is a separate subsystem that performs a specific operation (integer multiplication for instance). It consists of a datapath, control and interface. Typically it is interfaced with a start signal input and ready output (as well as data if required).
- The ALU is restricted to operations that can be performed quickly (hence the need for functional units).

- It's possible for the pipeline to issue a calculation to a functional unit, and then proceed to perform another instruction. It enables instructions to be performed in parallel and also for pipelines to get past a 'difficult' (time consuming) instruction.
- In superscalar parallelism, a processor uses multiple functional units to keep the pipeline going closer to full speed (obtaining parallelism among calculations). Programs are not restricted to the specific iterations over arrays supported by vector machines.
- Key concern is allowing as much parallelism as possible while maintaining correctness.
  - Centralised Control - The control algorithm analyses the stream of instructions and schedules the functional units to perform them.
  - Distributed control - Control is distributed among the subsystems (pipeline, functional units, signals etc).
- Tomasulo paper (refer)
- 

## Lecture 15 (Incomplete)

- Fast operations are done in the ALU, slow operations in dedicated functional units.
- Multiple functional units can operate in parallel.
- The functional units have two operands (sink & source) which receive data and a result buffer.
- When an operation reaches pipeline stage for executing the operation, the pipeline stalls, and the control unit performs the iteration using special registers in the main datapath.
- A functional unit has an interface (registers to receive operands from pipeline) and working registers (with combinational logic and control). In initial versions there was only 1 set of registers.

## Lecture 16

- Many different types of memory from small and fast to large and slow.
- The reason for the correlation between speed and size is because larger memory requires more address bits (and hence more time to decode the addresses).
- RAM is slower than registers because of the extra gate delays in decoding the address & the physical space that memory takes up (registers are directly located on the processor).
- The M1 only allows 1 clock cycle for a memory operation which is completely unrealistic (control algorithm must wait).
- The M1 control algorithm proceeds by placing whatever value is on the memory bus into the destination register (which will probably be incorrect).
- Interleaved memory is a solution to this problem. (used in the IBM 360/91).
- The memory consists of  $s^k$  individual memory banks, which is ordered sequentially.
- Many programs fetch or store a sequence of memory addresses. These can be done concurrently since they refer to distinct memory banks.

- Reduce random access by moving commonly accessed data into the faster cache.
- Cache mimics the most commonly used words in main memory but runs at full processor speed.
- Datapath and Control must determine if a value is in cache or not VERY quickly. Extra registers are needed to do this.
- Each location in the cache contains a word of data and a tag (the actual address of the data in memory). Which allows the machine to identify if an arbitrary address refers to data already in the cache, and if so, where it is.
- If data is in the cache it is a cache hit, else it is a cache miss (where the accessed value is then placed into the cache to prevent the same miss).
- A cache line is a sequence of data from memory in the same cache 'slot'. The greater the line the more likely a hit, but also the slower the access.
- For cache searching the most general approach is an associative search (each line has a comparator which check address against tag, this is done in parallel).
- The quickest approach is direct mapped cache which requires each address to go into a specific cache entry.
- Direct Mapped Cache - each word has a cache location, it contains data word, a tag, valid bit.
- Maps to a cache space based on its lower order k bits.
- If a value has 'valid' set then it contains the value of the word value.
- This is a bit inflexible. An alternative is to allow the value to do in any cache slot. The hardware must search all the tags of possible cache locations.
- Associative/Content-Addressable Memory. A field value is set and the memory return the rest of the word that matches that field is returned.
- Each location has logic as well. Values are compared in parallel, a tree of or gates can determine whether a match exists in logarithmic time. If there are multiple matches a tree circuit can determine a unique responder in logarithmic time.
- Compromise, set associative cache. A limited set of slots in cache that a value can be placed in.
- 2 approaches to memory stores, write-through (cache and memory values updates together), write-back (cache is allowed to become inconsistent with memory. Performance may be improved when this approach is taken).
- Virtual Memory is the like the difference between cache and memory elevated to the difference between memory and disk.
- A page table specifies where each part of the virtual address space actually is.
- Virtual memory offers a large address space that can (almost) be accessed at the speed of primary memory.
- Additional features include segmentation and protection.
- Pages are like cache lines in cache memory. Data isn't kept as words but rather sequences of words, typically 2Kb or 4Kb large. Either a page is in main memory or it is on the disk.


- A frame is a block of locations big enough to hold a page. When a page is in use it is kept in a frame.
- Page number and offset. The page number is the relative address of the page within the virtual address space, the offset is the specific location of the information relative to the page.
- Page table. A page table is an array (one entry per page) indexed by the page number. It contains entries 'resident' and 'frame address' which allow the system to determine if and where the page is in main memory. Other fields may exist depending on the system.
- Dynamic Address Translation (DAT). Effective Address  $\rightarrow$  Physical address in main memory OR Physical address on disk.
- Look over DAT algorithm.
- Dirty bit in a page table means that the page doesn't have to be written back to disk unless it has changed.
- A common page replacement algorithm is the 'least recently used' (LRU).
- Translation Lookaside Buffers are a special cache for the most active entries in the page table. It will usually have a reasonably large number of entries. Associative memory is used for TLBs.

## Lecture 17

- An interrupt is a jump (to the OS), initiated by an event other than a jump instruction.
- They are used to prevent user programs executing prohibited instructions, allow programs to access OS services, foundation for concurrent processes and threads.
- Interrupts are implemented as part of the control algorithm. If an interrupt request signal is 0 then the control algorithm functions as per usual. If 1 then the current pc is saved, and the updated to the address of the interrupt handler.
- The pc can be saved in a special register, into memory or onto a stack. All of these approaches have to deal with the possibility of other interrupts (and where to move the pc), what if there is a stack overflow.
- The state also needs to be saved. It can either be done by the interrupt handler (in software), which is flexible. Or be saved in the control algorithm (OS doesn't need to do it), which is simpler and reliable.
- As state could potentially be lost if an interrupt occurs whilst saving state. A signal 'enable' is added which only allows interrupts to take place when this signal is 1 (normally 0).
- If an interrupt is missed it must wait.
- Some events are 'real time' and there is no deadline by which they must be processed. OS should strive to save state quickly and enable interrupts.
- Memory protection is usually done with base and limit registers, where if the EA is  $> \text{lr}$  or  $< \text{br}$  then the operating system can decide what to do with the process. Most computers include this protection as part of the dynamic address translation algorithm in virtual memory.
- Some 'privileged' instructions can perform potentially dangerous operations (like altering the  $\text{lr}/\text{br}$  or I/O). These can only be executed by the OS.



- To check if it is a user or OS instruction a system state flip-flop is set. This is updated each time an instruction is received.
- Moore's law (see earlier). More complications with additional transistors. More transistors but not the same speed increase as 1990s.
- Techniques to notice speed increase (can generally only be applied once) include: Increased word size, on-chip cache memory, pipelining, superscalar and multiple threads. We are running out of speed increasing techniques.
- Trending towards large scale asynchronous circuits (though small circuits are synchronous, I/O is inherently asynchronous. ie Processors cannot insist input is given on its own clock).
- Clock distribution is much more difficult with larger circuits, it simply isn't practical given transistor trends.
- Use separate Intellectual Property (IP) cores, made by different vendors but installed on the same chip. Intercommunication is done via 'network on a chip'
- If a signal is produced by a circuit on clock 1, and received by a flip flop on clock 2 it is asynchronous.
- The signal may not have enough time to settle down (and be incorrect) or may be in transition (metastable).
- Metastability is the area between a 1 and a 0. Signals snap to either 1 (even if there is minor degradation). Flip flops typically become metastable with large degradation or transitions. They will eventually snap to a 1/0 after a random (but unbounded) time.
- A signal has to be converted to make it valid in the receivers clock domain (synchronization).



- Synchronizer doesn't provide a guarantee of metastability but reduces its likelihood. There must be a trade off between an acceptable probability of failure and speed/power consumption of the synchronizer.
- GALS - Globally Asynchronous Locally Synchronous. (As it says on the tin).